

# 1ST EVALUATION


```
In [134]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
dataset = pd.read_csv('C:\\Users\\kama1\\OneDrive\\Documents\\Machine Learning
```

```
In [135]: dataset
```

```
Out[135]:
```

	Creditability	Account Balance	Duration of Credit (month)	Payment Status of Previous Credit	Purpose	Credit Amount	Value Savings/Stocks	Length of current employment
0	1	1	18	4	2	1049	1	2
1	1	1	9	4	0	2799	1	3
2	1	2	12	2	9	841	2	4
3	1	1	12	4	0	2122	1	3
4	1	1	12	4	0	2171	1	3
...	...	...	...	...	...	...	...	...
995	0	1	24	2	3	1987	1	3
996	0	1	24	2	0	2303	1	5
997	0	4	21	4	0	12680	5	5
998	0	2	12	2	3	6468	5	1
999	0	1	30	2	2	6350	5	5

1000 rows × 21 columns



```
In [136]: missing = dataset.isnull().sum()
```

In [137]: missing


```
Out[137]: Creditability      0
Account Balance            0
Duration of Credit (month) 0
Payment Status of Previous Credit 0
Purpose                    0
Credit Amount              0
Value Savings/Stocks       0
Length of current employment 0
Instalment per cent        0
Sex & Marital Status       0
Guarantors                 0
Duration in Current address 0
Most valuable available asset 0
Age (years)                0
Concurrent Credits         0
Type of apartment          0
No of Credits at this Bank  0
Occupation                 0
No of dependents           0
Telephone                  0
Foreign Worker             0
dtype: int64
```

In [138]: dataset.describe()

```
Out[138]:
```

	Creditability	Account Balance	Duration of Credit (month)	Payment Status of Previous Credit	Purpose	Credit Amount	Value Savings/Stocks
<b>count</b>	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.00
<b>mean</b>	0.700000	2.577000	20.903000	2.54500	2.828000	3271.24800	2.10
<b>std</b>	0.458487	1.257638	12.058814	1.08312	2.744439	2822.75176	1.58
<b>min</b>	0.000000	1.000000	4.000000	0.00000	0.000000	250.00000	1.00
<b>25%</b>	0.000000	1.000000	12.000000	2.00000	1.000000	1365.50000	1.00
<b>50%</b>	1.000000	2.000000	18.000000	2.00000	2.000000	2319.50000	1.00
<b>75%</b>	1.000000	4.000000	24.000000	4.00000	3.000000	3972.25000	3.00
<b>max</b>	1.000000	4.000000	72.000000	4.00000	10.000000	18424.00000	5.00

8 rows × 21 columns

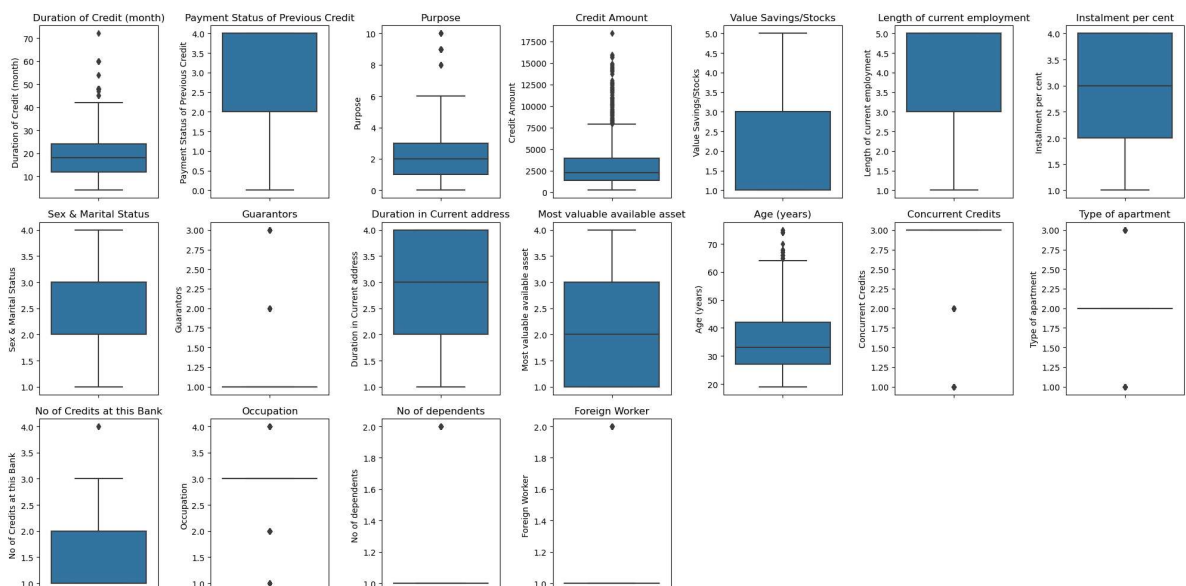


```
In [139]: print(dataset.dtypes)
```

```
Creditability          int64
Account Balance        int64
Duration of Credit (month)  int64
Payment Status of Previous Credit  int64
Purpose                int64
Credit Amount          int64
Value Savings/Stocks   int64
Length of current employment  int64
Instalment per cent    int64
Sex & Marital Status    int64
Guarantors             int64
Duration in Current address  int64
Most valuable available asset  int64
Age (years)            int64
Concurrent Credits      int64
Type of apartment       int64
No of Credits at this Bank  int64
Occupation              int64
No of dependents        int64
Telephone              int64
Foreign Worker          int64
dtype: object
```

```
In [140]: box_plot = dataset.drop(columns=['Creditability', 'Account Balance', 'Telephone'])
```

```
In [141]: plt.figure(figsize=(20, 10))
for i, col in enumerate(box_plot):
    plt.subplot(3, 7, i + 1)
    sns.boxplot(y=dataset[col])
    plt.title(col)
plt.tight_layout()
plt.show()
```



```
In [142]: attributes = ['Creditability', 'Account Balance', 'Duration of Credit (month)',
                        'Length of current employment', 'Instalment per cent', 'Sex & Marital S',
                        'Type of apartment', 'No of Credits at this Bank', 'Occupation', 'No of']
```

```
In [143]: import pandas as pd
          from scipy.stats import zscore

          z_scores = dataset[attributes].apply(zscore)
          threshold = 3

          outliers_mask = abs(z_scores) > threshold
          outliers = dataset[outliers_mask]

          sum_of_outliers = outliers_mask.sum()

          print("Sum of outliers in each column:")
          print(sum_of_outliers)
```

```
Sum of outliers in each column:
Creditability                0
Account Balance              0
Duration of Credit (month)   14
Payment Status of Previous Credit  0
Purpose                      0
Credit Amount               25
Value Savings/Stocks        0
Length of current employment  0
Instalment per cent         0
Sex & Marital Status         0
Duration in Current address  0
Most valuable available asset 0
Age (years)                  7
Concurrent Credits          0
Type of apartment           0
No of Credits at this Bank   6
Occupation                  0
No of dependents            0
Telephone                   0
Foreign Worker              37
dtype: int64
```

In [175]: `#outlier_columns = ['Duration of Credit (month)', 'Credit Amount', 'Age (years)`

```
#def remove_outliers(df, columns):
#    for column in columns:
#        mean = dataset[column].mean()
#        std = dataset[column].std()
#        threshold = 3
#        df = df[(df[column] - mean).abs() < threshold * std]
#    return df

#df = remove_outliers(dataset, outlier_columns)
```

In [144]: `for column, value in outliers.items():
 dataset.loc[dataset[column] == value, column] = dataset[column].mean()

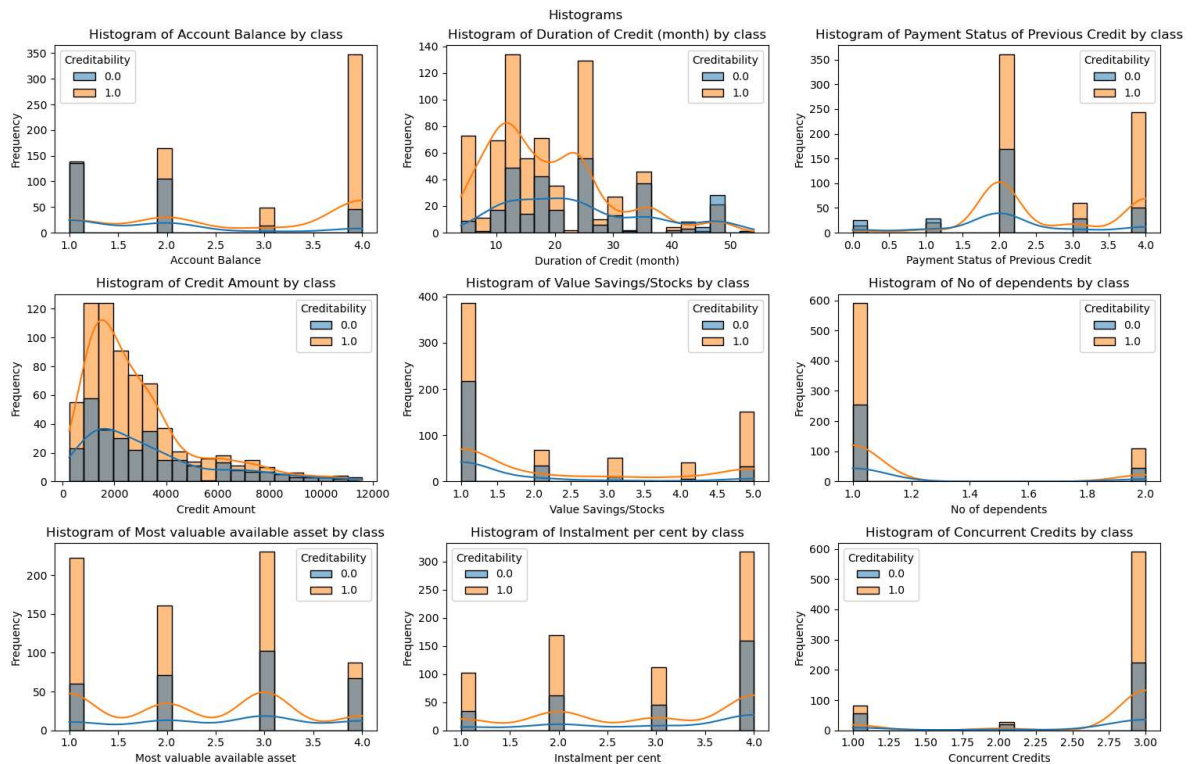
print(dataset)`

	Creditability	Account Balance	Duration of Credit (month)	\
0	1.0	1.0	18.0	
1	1.0	1.0	9.0	
2	1.0	2.0	12.0	
3	1.0	1.0	12.0	
4	1.0	1.0	12.0	
..	...	...	...	
995	0.0	1.0	24.0	
996	0.0	1.0	24.0	
997	0.0	4.0	21.0	
998	0.0	2.0	12.0	
999	0.0	1.0	30.0	

	Payment Status of Previous Credit	Purpose	Credit Amount	\
0	4.0	2.0	1049.000	
1	4.0	0.0	2799.000	
2	2.0	9.0	841.000	
3	4.0	0.0	2122.000	
4	4.0	0.0	2171.000	

```
In [145]: plt.figure(figsize=(15, 10))
for i, col in enumerate(['Account Balance', 'Duration of Credit (month)', 'Pay
plt.subplot(3, 3, i + 1)
sns.histplot(data=dataset, x=col, hue='Creditability', kde=True, bins=20)
plt.title(f'Histogram of {col} by class')
plt.xlabel(col)
plt.ylabel('Frequency')
plt.suptitle('Histograms')
plt.tight_layout()
plt.show()
```



```
In [146]: numeric_cols = dataset.select_dtypes(include=[np.number]).columns
numeric_cols
```

```
Out[146]: Index(['Creditability', 'Account Balance', 'Duration of Credit (month)',
'Payment Status of Previous Credit', 'Purpose', 'Credit Amount',
'Value Savings/Stocks', 'Length of current employment',
'Instalment per cent', 'Sex & Marital Status', 'Guarantors',
'Duration in Current address', 'Most valuable available asset',
'Age (years)', 'Concurrent Credits', 'Type of apartment',
'No of Credits at this Bank', 'Occupation', 'No of dependents',
'Telephone', 'Foreign Worker'],
dtype='object')
```

```
In [147]: attributes = ['Creditability', 'Account Balance', 'Duration of Credit (month)',
                        'Length of current employment', 'Instalment per cent', 'Sex & Marital S
                        'Type of apartment', 'No of Credits at this Bank', 'Occupation', 'No of

correlation_matrix = dataset[attributes].corr()


print("Correlation Matrix:")
print(correlation_matrix)
```

Correlation Matrix:

	Creditability	Account Balance \
Creditability	1.000000	0.350847
Account Balance	0.350847	1.000000
Duration of Credit (month)	-0.210833	-0.077712
Payment Status of Previous Credit Purpose	0.228785	0.192191
Credit Amount	-0.017979	0.028783
Value Savings/Stocks	-0.095433	-0.010968
Length of current employment	0.178943	0.222867
Instalment per cent	0.116002	0.106339
Sex & Marital Status	-0.072404	-0.005280
Duration in Current address	0.088184	0.043261
Most valuable available asset	-0.002967	-0.042234
Age (years)	-0.142612	-0.032260
Concurrent Credits	0.086792	0.064106
Type of apartment	0.109844	0.068274
No of Credits at this Bank	0.018119	0.023335
Occupation	0.050891	0.086676
No of Credits at this Bank	-0.032735	0.040663
No of Credits at this Bank	0.000000	0.000000

```
In [148]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=1)
plt.title('Correlation Matrix Heatmap')
plt.show()
```





```
In [149]: correlation_matrix = dataset.corr()
correlation_matrix
```

Out[149]:

	Creditability	Account Balance	Duration of Credit (month)	Payment Status of Previous Credit	Purpose	Credit Amount	Value Savings/Stock
<b>Creditability</b>	1.000000	0.350847	-0.210833	0.228785	-0.017979	-0.095433	0.17894
<b>Account Balance</b>	0.350847	1.000000	-0.077712	0.192191	0.028783	-0.010968	0.22286
<b>Duration of Credit (month)</b>	-0.210833	-0.077712	1.000000	-0.074935	0.148448	0.578733	0.04184
<b>Payment Status of Previous Credit</b>	0.228785	0.192191	-0.074935	1.000000	-0.090336	-0.006512	0.03905
<b>Purpose</b>	-0.017979	0.028783	0.148448	-0.090336	1.000000	0.008592	-0.01868
<b>Credit Amount</b>	-0.095433	-0.010968	0.578733	-0.006512	0.008592	1.000000	0.07138
<b>Value Savings/Stocks</b>	0.178943	0.222867	0.041845	0.039058	-0.018684	0.071381	1.00000
<b>Length of current employment</b>	0.116002	0.106339	0.043461	0.138225	0.016013	-0.006342	0.12095
<b>Instalment per cent</b>	-0.072404	-0.005280	0.090636	0.044375	0.048369	-0.281638	0.02199
<b>Sex &amp; Marital Status</b>	0.088184	0.043261	0.001848	0.042171	0.000157	-0.017162	0.01734
<b>Guarantors</b>	0.025137	-0.127737	-0.018721	-0.040676	-0.017607	-0.027517	-0.10506
<b>Duration in Current address</b>	-0.002967	-0.042234	0.021376	0.063198	-0.038221	0.020600	0.09142
<b>Most valuable available asset</b>	-0.142612	-0.032260	0.293408	-0.053777	0.010966	0.275370	0.01894
<b>Age (years)</b>	0.086792	0.064106	-0.027445	0.148430	-0.001416	-0.011265	0.09497
<b>Concurrent Credits</b>	0.109844	0.068274	-0.071741	0.159957	-0.100230	-0.038534	0.00190
<b>Type of apartment</b>	0.018119	0.023335	0.133234	0.061428	0.013495	0.101348	0.00664
<b>No of Credits at this Bank</b>	0.050891	0.086676	0.003613	0.449865	0.049828	0.073000	-0.03332
<b>Occupation</b>	-0.032735	0.040663	0.195234	0.010350	0.008085	0.254145	0.01170
<b>No of dependents</b>	0.003015	-0.014145	-0.014072	0.011550	-0.032577	0.049354	0.02751
<b>Telephone</b>	0.036466	0.066296	0.163410	0.052370	0.078371	0.224689	0.08720
<b>Foreign Worker</b>	0.082079	-0.035187	-0.136772	0.028554	-0.113244	-0.051240	0.01045

21 rows × 21 columns

## 2ND EVALUATION

```
In [165]: x = dataset[['Account Balance', 'Duration of Credit (month)', 'Payment Status of  
            'Value Savings/Stocks', 'Age (years)', 'Occupation', 'Instalment per cent',  
            'Most valuable available asset', 'Concurrent Credits', 'No of Credits at th  
            'No of dependents', 'Foreign Worker']]  
y = dataset['Creditability']
```

```
In [166]: from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [167]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
x_train_scaled = scaler.fit_transform(x_train)  
x_test_scaled = scaler.fit_transform(x_test)
```

## Naive Bayes Model

```
In [168]: from sklearn.naive_bayes import GaussianNB  
from sklearn import metrics  
gnb = GaussianNB()
```

```
In [169]: gnb.fit(x_train_scaled, y_train)  
y_pred = gnb.predict(x_test_scaled)  
nb_accuracy = accuracy_score(y_test, y_pred)
```

## Decision tree classifier

```
In [170]: from sklearn.tree import DecisionTreeClassifier  
dr = DecisionTreeClassifier()  
dr.fit(x_train_scaled, y_train)  
predict = dr.predict(x_test_scaled)  
dr_accuracy = accuracy_score(y_test, predict)
```

## Logistic Regression Model

```
In [171]: from sklearn.linear_model import LogisticRegression  
lr = LogisticRegression()  
lr.fit(x_train_scaled, y_train)  
predict_lr = lr.predict(x_test_scaled)  
lr_accuracy = accuracy_score(y_test, predict_lr)
```

## Comparing accuracy

```
In [172]: print("Accuracy of naive bayes :", nb_accuracy)
          print("Accuracy of Decision tree: ", dr_accuracy)
          print("Accuracy based on logistic regression model: ", lr_accuracy)
```

```
Accuracy of naive bayes : 0.73
Accuracy of Decision tree: 0.675
Accuracy based on logistic regression model: 0.72
```

## Comparing recall

```
In [173]: from sklearn.metrics import recall_score

          recall_naivebayes = recall_score(y_test, y_pred)
          recall_dr = recall_score(y_test, predict)
          recall_lr = recall_score(y_test, predict_lr)

          print("Recall for naive bayes:", recall_naivebayes)
          print("Recall for Decision Tree:", recall_dr)
          print("Recall for Logistic Regression :", recall_lr)
```

```
Recall for naive bayes: 0.7867647058823529
Recall for Decision Tree: 0.8161764705882353
Recall for Logistic Regression : 0.875
```

## Comparing f1 score

```
In [174]: from sklearn.metrics import f1_score

          f1_naivebayes = f1_score(y_test, y_pred)
          f1_dr = f1_score(y_test, predict)
          f1_lr = f1_score(y_test, predict_lr)

          print("f1 score for naive bayes :", f1_naivebayes)
          print("f1 score for decision tree :", f1_dr)
          print("f1 score for logistic regression :", f1_lr)
```

```
f1 score for naive bayes : 0.7985074626865671
f1 score for decision tree : 0.7735191637630662
f1 score for logistic regression : 0.8095238095238096
```

## COMPARING RESULTS

**1) The decision tree model performs the worst among the three models as it has an accuracy of 0.71, which is the lowest. Its recall and F1 score are also lower than those of the other models. This**

**indicates that the decision tree model may not be the best choice for this dataset compared to logistic regression and naive Bayes.**

**2) The naive Bayes model might be an appropriate choice for this this dataset with an accuracy of 0.75**

**3) The logistic regression model outperforms both the naive Bayes and decision tree models in terms of accuracy, recall, and F1 score. It has the highest accuracy of 0.76, recall - 0.90 and f1 score - 0.84**

**From the above observations, a logistic regression model is the most suitable model for this dataset**

```
In [160]: import pandas as pd
filtered_data = dataset[dataset['Creditability'] == 1]
```

```
In [161]: from sklearn.cluster import KMeans

wcss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(filtered_data)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss)
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.title('Elbow Method')
plt.show()
```

```
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:87
0: FutureWarning: The default value of `n_init` will change from 10 to 'au
to' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:138
2: UserWarning: KMeans is known to have a memory leak on Windows with MKL,
when there are less chunks than available threads. You can avoid it by set
ting the environment variable OMP_NUM_THREADS=3.
  warnings.warn(
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:87
0: FutureWarning: The default value of `n_init` will change from 10 to 'au
to' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:138
2: UserWarning: KMeans is known to have a memory leak on Windows with MKL,
when there are less chunks than available threads. You can avoid it by set
ting the environment variable OMP_NUM_THREADS=3.
  warnings.warn(
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:87
```

```
In [162]: k = 2
kmeans = KMeans(n_clusters=k, random_state=42)
kmeans.fit(filtered_data)
```

```
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Program Files\Anaconda\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=3.
  warnings.warn(
```

```
Out[162]: KMeans(n_clusters=2, random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [163]: from sklearn.metrics import silhouette_score

silhouette_avg = silhouette_score(filtered_data, kmeans.labels_)

# Print or interpret the score
print("Silhouette Score (k =", k, "):", silhouette_avg)
```

```
Silhouette Score (k = 2 ): 0.7010705450093719
```

## CONCLUSIONS

### Before treating outliers

Accuracy of logistic regression before treating outliers- 0.76

recall of logistic regression before treating outliers - 0.90

f1 score of logistic regression before treating outliers - 0.84

### After replacing outliers with mean

Accuracy based on logistic regression model: 0.72

Recall for Logistic Regression : 0.875

f1 score for logistic regression : 0.8095238095238096

## After removing outliers

Accuracy based on logistic regression model: 0.7248677248677249

Recall for Logistic Regression : 0.889763779527559

f1 score for logistic regression : 0.8129496402877697

1)The accuracy of the logistic regression model decreased slightly after replacing outliers with the mean and after removing outliers compared to before treating outliers. This suggests that treating outliers did not significantly improve the overall accuracy of the model.

However, the recall of the logistic regression model improved after both outlier treatment methods. This indicates that the model became better at correctly identifying positive cases (credible applicants in this context) after outlier treatment.

The F1 score, which is a measure of a model's accuracy, also showed improvement after outlier treatment, indicating a better balance between precision and recall.

**This suggests that the outliers in the dataset are not due to errors but rather because the dataset includes values at a higher range that are important for the prediction model.**

2)K-means clustering is a method used to partition a dataset into groups (clusters) based on similarities in the data.

The Silhouette Score is a measure of how similar an object is to its own cluster compared to other clusters. A score close to 1 indicates that the object is well matched to its own cluster and poorly matched to neighboring clusters, suggesting a good clustering. In this case, a score of 0.701 suggests that the clustering with 2 clusters is reasonably good

## RECOMMENDATIONS

These are the features which will have a greater or a higher impact on the target variable

Payment Status of Previous Credit

Credit Amount

Duration of Credit (month)

Account Balance

Value Savings/Stocks

No of Credits at this Bank

No of dependents

## Foreign Worker

These are the features which can be removed as they do not affect the target variable to a greater extent-

Purpose

Sex & Marital Status

Type of apartment

Telephone

In [ ]: